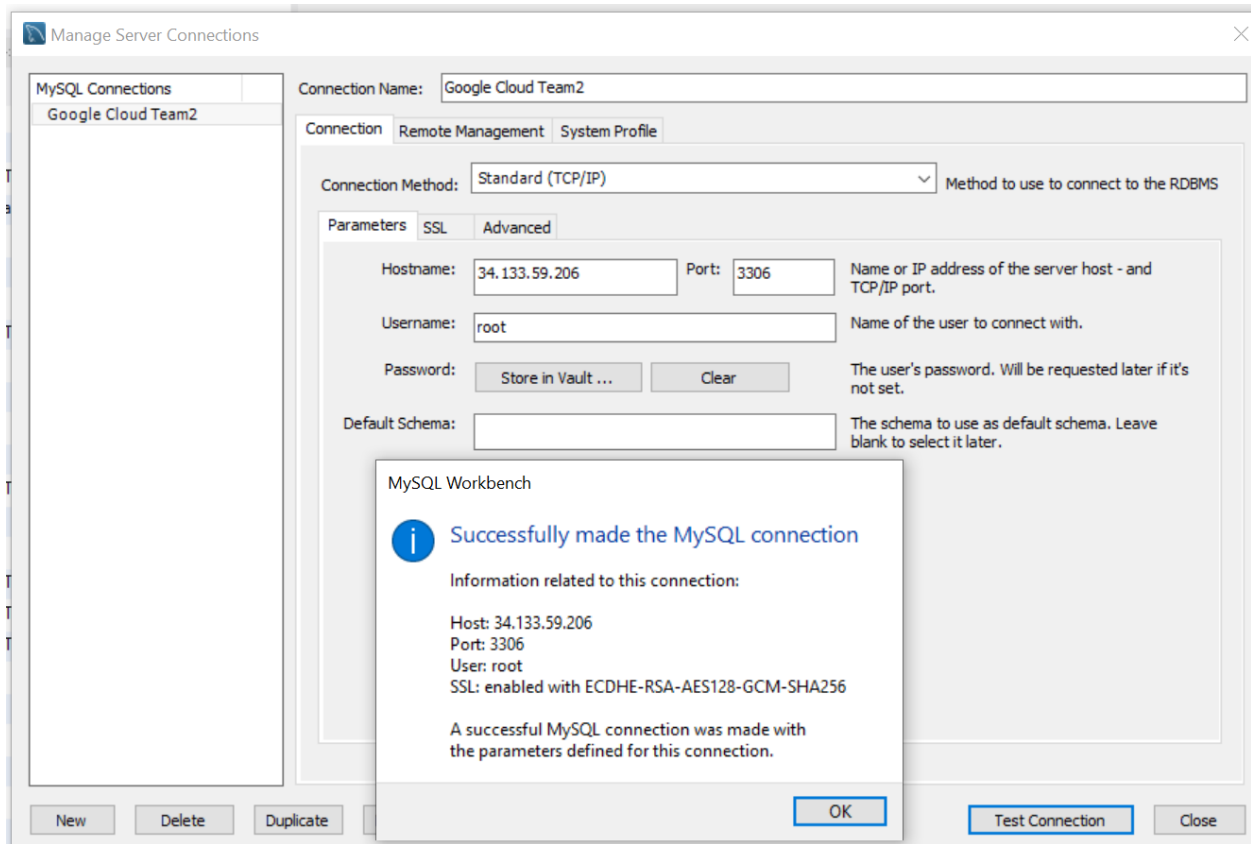


A. Create a markdown or pdf called “Database Design” in the doc folder

GCP Connection:



DDL for Table Creation:

```
CREATE TABLE Athletes
(
    Name VARCHAR(512) NOT NULL,
    NOC VARCHAR(512) NOT NULL,
    Discipline VARCHAR(512) NOT NULL,
    PRIMARY KEY (Name, Discipline)
);

CREATE TABLE Coaches
(
    Name VARCHAR(512),
    NOC VARCHAR(512),
    Discipline VARCHAR(512),
    Event VARCHAR(512),
    PRIMARY KEY (Name, Event)
```

```

);
CREATE TABLE Medals
(
    NOC VARCHAR(512),
    RankOverall INT,
    Gold INT,
    Silver INT,
    Bronze INT,
    Total INT,
    RankByNumberOfMedals INT
    PRIMARY KEY (NOC)
);
CREATE TABLE MedalsByAthlete
(
    athlete_name VARCHAR(512),
    event VARCHAR(512),
    NOC VARCHAR(512),
    discipline VARCHAR(512),
    medal_type VARCHAR(512),
    medal_date VARCHAR(512),
    PRIMARY KEY (athlete_name, event)
);
CREATE TABLE Officials
(
    name VARCHAR(512),
    gender VARCHAR(512),
    country VARCHAR(512),
    discipline VARCHAR(512),
    role VARCHAR(512),
    PRIMARY KEY (name)
);

```

Number of Entries per table:

Athletes:

1 • `select count(*) from Athletes;`
2

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

count(*)
11077

Coaches:

1 • `select count(*) from Coaches;`
2

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

count(*)
380

Medals:

1 • `select count(*) from Medals;`
2

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

count(*)
93

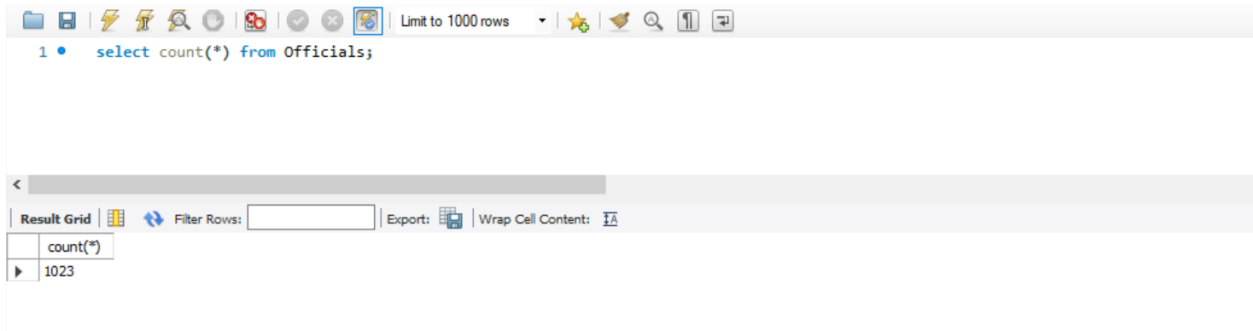
MedalsByAthlete:

1 • `select count(*) from MedalsByAthlete;`
2

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

count(*)
2401

Officials:



The screenshot shows a SQL query editor with a toolbar at the top. The query entered is `select count(*) from Officials;`. Below the query editor, the 'Result Grid' is displayed, showing a single row with the column `count(*)` and the value `1023`. The interface includes a 'Filter Rows' input field, an 'Export' button, and a 'Wrap Cell Content' checkbox.

count(*)
1023

#1 Advanced SQL Query

```

SELECT distinct
  Athletes.Discipline, Coaches.Name
FROM
  Coaches
  INNER JOIN
    Athletes ON Coaches.NOC = Athletes.NOC
    AND Coaches.Discipline = Athletes.Discipline
  INNER JOIN
    MedalsByAthlete ON Athletes.Name = MedalsByAthlete.Athlete_Name
WHERE
  MedalsByAthlete.medal_type = 'Gold Medal'
LIMIT 15;

;

```

Image:

```

1 • SELECT distinct
2     Athletes.Discipline, Coaches.Name
3 FROM
4     Coaches
5     INNER JOIN
6     Athletes ON Coaches.NOC = Athletes.NOC
7     AND Coaches.Discipline = Athletes.Discipline
8     INNER JOIN
9     MedalsByAthlete ON Athletes.Name = MedalsByAthlete.Athlete_Name
10 WHERE
11     MedalsByAthlete.medal_type = 'Gold Medal'
12 LIMIT 15;
13

```

Discipline	Name
Hockey	ANNAN Alyson
Rugby Sevens	BABER Gareth Colin
Rugby Sevens	BUNTING Allan
Hockey	CALDAS Max
Baseball/Softball	CASAREZ Mary Lucy
Hockey	COMMENS Adam
Artistic Swimming	DANCHENKO Tatiana
Hockey	de VOS Robbert-Jan
Basketball	ENGELLAND Chip

Result 10 x Read Only

Justification: There are many countries whose coaches do not get recognition for contributing to a nation that wins medals in these competitive games. This advanced query uses INNER JOIN and a subquery. It finds all of the coaches and connects it to the medals table so the output is only a list of coaches whose athletes have won a gold medal.

#2 Advanced SQL Query

Explain Analyze

```

SELECT a.Name AS LosingAthlete, mba.athlete_name AS WinningAthlete, a.discipline
FROM Athletes a
LEFT JOIN (
    SELECT DISTINCT mba.discipline, mba.athlete_name
    FROM MedalsByAthlete mba
    WHERE mba.medal_type = 'Gold Medal'
) mba ON a.Discipline = mba.discipline
WHERE a.Name NOT IN (
    SELECT Distinct athlete_name
    FROM MedalsByAthlete
    WHERE medal_type IN ('Gold Medal', 'Silver Medal',

```

'Bronze Medal')
);

Image:

The image shows a SQL IDE interface. The top pane contains a SQL query. The bottom pane shows the results of the query in a table format.

```
1  -- Explain Analyze
2  • SELECT  a.Name AS LosingAthlete, mba.athlete_name AS WinningAthlete, a.discipline
3  FROM Athletes a
4  LEFT JOIN (
5      SELECT DISTINCT mba.discipline, mba.athlete_name
6      FROM MedalsByAthlete mba
7      WHERE mba.medal_type = 'Gold Medal'
8  ) mba ON a.Discipline = mba.discipline
9  WHERE a.Name NOT IN (
10     SELECT Distinct athlete_name
11     FROM MedalsByAthlete
12     WHERE medal_type IN ('Gold Medal', 'Silver Medal',
13                          'Bronze Medal')
14 )
15 LIMIT 15;
```

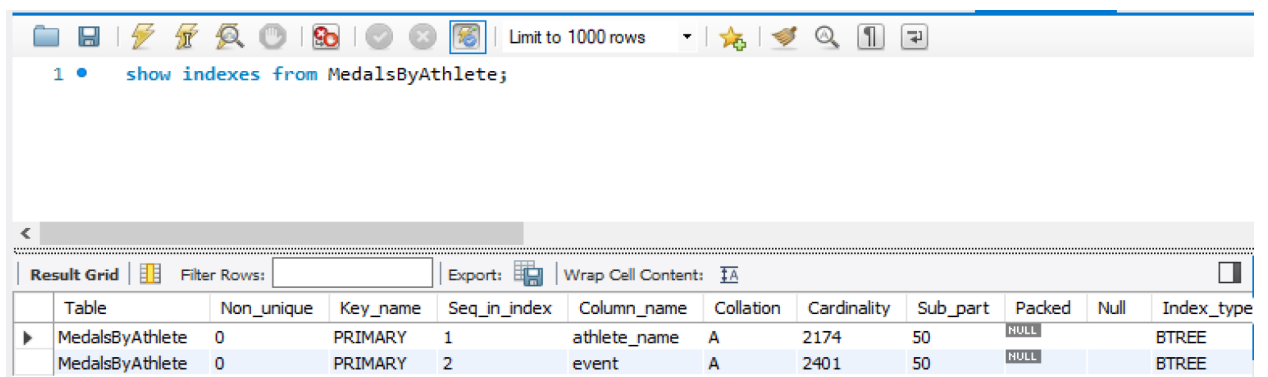
LosingAthlete	WinningAthlete	discipline
AALERUD Katrine	van VLEUTEN Annemiek	Cycling Road
AALERUD Katrine	ROGLIC Primoz	Cycling Road
AALERUD Katrine	KIESENHOFER Anna	Cycling Road
AALERUD Katrine	CARAPAZ Richard	Cycling Road
ABAD Nestor	ZOU Jingyuan	Artistic Gymnastics
ABAD Nestor	WHITLOCK Max	Artistic Gymnastics
ABAD Nestor	URAZOVA Vladislava	Artistic Gymnastics
ABAD Nestor	SHIN Jeahwan	Artistic Gymnastics
ABAD Nestor	NAGORNYI Nikita	Artistic Gymnastics
ABAD Nestor	MELNIKOVA Angelina	Artistic Gymnastics
ABAD Nestor	LIU Yang	Artistic Gymnastics
ABAD Nestor	LISTUNOVA Viktoriia	Artistic Gymnastics
ABAD Nestor	LEE Sunisa	Artistic Gymnastics
ABAD Nestor	HASHIMOTO Daiki	Artistic Gymnastics
ABAD Nestor	GUAN Chenchen	Artistic Gymnastics

Justification: For each athlete that didn't win a medal, we want to display the athlete that won the gold medal for that discipline.

[Indexing]

Advanced Query 1:

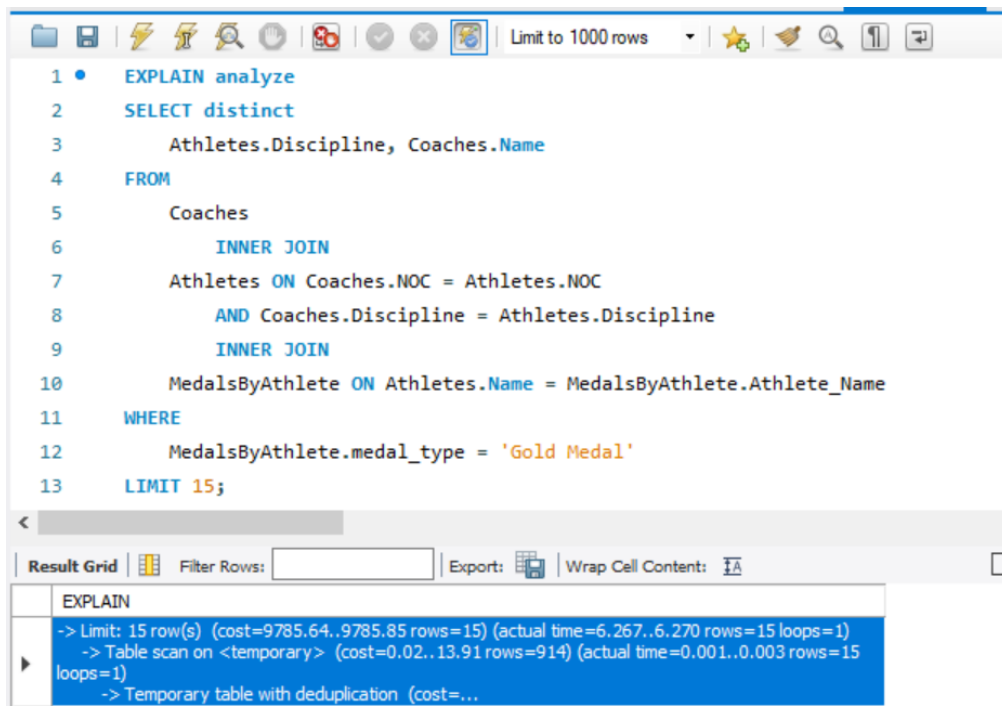
1. Indexes before:



```
1 • show indexes from MedalsByAthlete;
```

	Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type
▶	MedalsByAthlete	0	PRIMARY	1	athlete_name	A	2174	50	NULL		BTREE
	MedalsByAthlete	0	PRIMARY	2	event	A	2401	50	NULL		BTREE

Current runtime:



```
1 • EXPLAIN analyze
2 SELECT distinct
3     Athletes.Discipline, Coaches.Name
4 FROM
5     Coaches
6     INNER JOIN
7     Athletes ON Coaches.NOC = Athletes.NOC
8     AND Coaches.Discipline = Athletes.Discipline
9     INNER JOIN
10    MedalsByAthlete ON Athletes.Name = MedalsByAthlete.Athlete_Name
11 WHERE
12     MedalsByAthlete.medal_type = 'Gold Medal'
13 LIMIT 15;
```

	EXPLAIN
▶	-> Limit: 15 row(s) (cost=9785.64..9785.85 rows=15) (actual time=6.267..6.270 rows=15 loops=1) -> Table scan on <temporary> (cost=0.02..13.91 rows=914) (actual time=0.001..0.003 rows=15 loops=1) -> Temporary table with deduplication (cost=...

New Indexes:

Query 1 SQL File 1* **SQL File 2*** SQL File 3* SQL File 4* SQL File 5* SQL File 6*

1 • `alter table MedalsByAthlete add index (medal_type(50));`
 2 • `show indexes from MedalsByAthlete;`
 3

Result Grid Filter Rows: Export: Wrap Cell Content: **Result Grid**

	Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part
▶	MedalsByAthlete	0	PRIMARY	1	athlete_name	A	2174	50
	MedalsByAthlete	0	PRIMARY	2	event	A	2401	50
	MedalsByAthlete	1	medal_type	1	medal_type	A	3	50

New Runtime:

1 • `EXPLAIN analyze`
 2 `SELECT distinct`
 3 `Athletes.Discipline, Coaches.Name`
 4 `FROM`
 5 `Coaches`
 6 `INNER JOIN`
 7 `Athletes ON Coaches.NOC = Athletes.NOC`
 8 `AND Coaches.Discipline = Athletes.Discipline`
 9 `INNER JOIN`
 10 `MedalsByAthlete ON Athletes.Name = MedalsByAthlete.Athlete_Name`
 11 `WHERE`
 12 `MedalsByAthlete.medal_type = 'Gold Medal'`
 13 `LIMIT 15;`

Result Grid Filter Rows: Export: Wrap Cell Content: **Result Grid**

	EXPLAIN
▶	-> Limit: 15 row(s) (cost=357332.56..357332.74 rows=15) (actual time=12.514..12.517 rows=15 loops=1) -> Table scan on <temporary> (cost=0.01..373.48 rows=29678) (actual time=0.003..0.004 rows=15 loops=1) -> Temporary table with deduplicatio...

Result: Twice as slow

Default Indexing And Runtime = 6.275:

1 • `show indexes from Coaches;`
 2 `-- show indexes from MedalsByAthlete;`
 3

Result Grid Filter Rows: Export: Wrap Cell Content: **Result Grid**

	Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	I
▶	Coaches	0	PRIMARY	1	Name	A	374	50	NULL		B
	Coaches	0	PRIMARY	2	Event	A	380	50	NULL		B

Limit to 1000 rows

```

1 -- show indexes from Coaches;
2 • show indexes from MedalsByAthlete;
3

```

Result Grid

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	I
MedalsByAthlete	0	PRIMARY	1	athlete_name	A	2174	50		NULL	
MedalsByAthlete	0	PRIMARY	2	event	A	2401	50		NULL	

Adding index to Coaches.NOC:

Limit to 1000 rows

```

1 • alter table Coaches add index (NOC(50));
2 • show indexes from Coaches;
3

```

Result Grid

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	I
Coaches	0	PRIMARY	1	Name	A	374	50		NULL	B
Coaches	0	PRIMARY	2	Event	A	380	50		NULL	B
Coaches	1	NOC	1	NOC	A	61	50		NULL	YES B

Runtime:

The screenshot shows a SQL IDE window with a query editor and a result grid. The query is as follows:

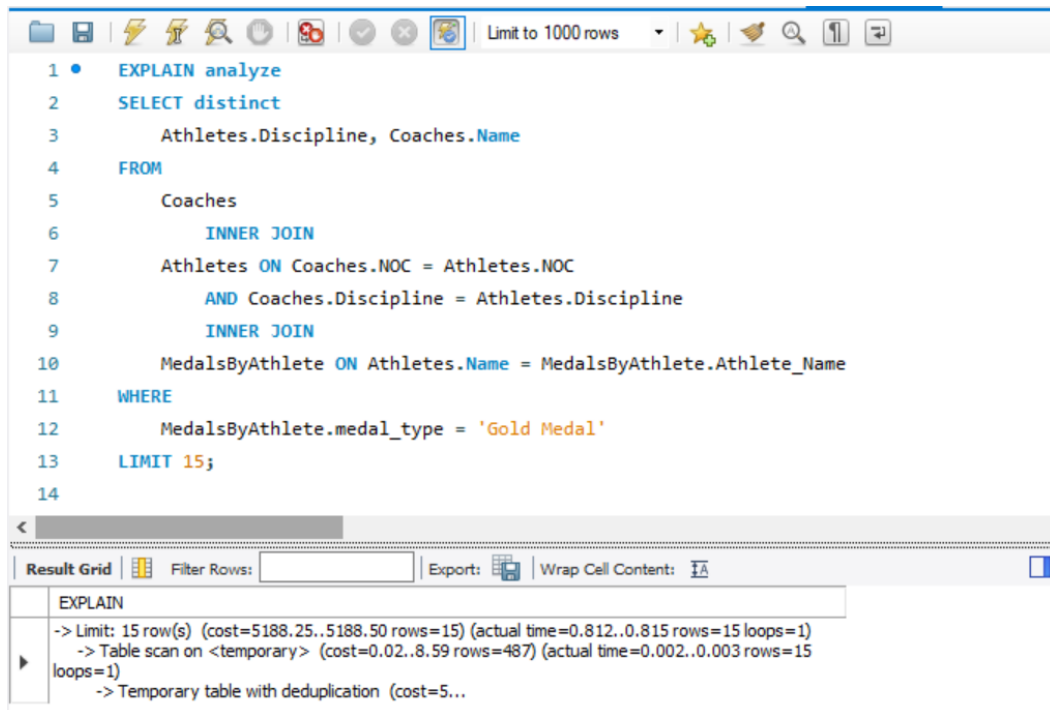
```
1 • EXPLAIN analyze
2 SELECT distinct
3     Athletes.Discipline, Coaches.Name
4 FROM
5     Coaches
6     INNER JOIN
7     Athletes ON Coaches.NOC = Athletes.NOC
8     AND Coaches.Discipline = Athletes.Discipline
9     INNER JOIN
10    MedalsByAthlete ON Athletes.Name = MedalsByAthlete.Athlete_Name
11 WHERE
12    MedalsByAthlete.medal_type = 'Gold Medal'
13 LIMIT 15;
14
```

The result grid shows the execution plan for the query:

EXPLAIN
-> Limit: 15 row(s) (cost=1812.92..1813.33 rows=15) (actual time=4.523..4.527 rows=15 loops=1)
-> Table scan on <temporary> (cost=0.03..4.36 rows=150) (actual time=0.002..0.004 rows=15 loops=1)
-> Temporary table with deduplication (cost=1...

Result: The original cost of the algorithm was 9785 with default indexing, but by indexing Coaches.NOC, it went down to 1812.

Indexing both MedalsByAthlete.medal_type and Coaches.NOC
Runtime:



```
1 • EXPLAIN analyze
2 SELECT distinct
3     Athletes.Discipline, Coaches.Name
4 FROM
5     Coaches
6     INNER JOIN
7     Athletes ON Coaches.NOC = Athletes.NOC
8     AND Coaches.Discipline = Athletes.Discipline
9     INNER JOIN
10    MedalsByAthlete ON Athletes.Name = MedalsByAthlete.Athlete_Name
11 WHERE
12     MedalsByAthlete.medal_type = 'Gold Medal'
13 LIMIT 15;
14
```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: [\[A\]](#)

EXPLAIN
-> Limit: 15 row(s) (cost=5188.25..5188.50 rows=15) (actual time=0.812..0.815 rows=15 loops=1)
-> Table scan on <temporary> (cost=0.02..8.59 rows=487) (actual time=0.002..0.003 rows=15 loops=1)
-> Temporary table with deduplication (cost=5...

Results: Faster than default indexing, but for some reason adding indexing MedalsByAthlete.medal_type makes the algorithm slower, even when paired with indexing on Coaches.NOC.

We will choose to add index to Coaches.NOC but not index MedalsByAthlete.medal_type as this will produce the lowest runtime.

2.

Advanced Query 2

Default Indexing runtime:

The screenshot displays a SQL IDE with a query editor and an execution plan window.

Query Editor:

```
1 • Explain Analyze
2 SELECT a.Name AS LosingAthlete, mba.athlete_name AS WinningAthlete, a.discipline
3 FROM Athletes a
4 LEFT JOIN (
5     SELECT DISTINCT mba.discipline, mba.athlete_name
6     FROM MedalsByAthlete mba
7     WHERE mba.medal_type = 'Gold Medal'
8 ) mba ON a.Discipline = mba.discipline
9 WHERE a.Name NOT IN (
10     SELECT Distinct athlete_name
11     FROM MedalsByAthlete
12     WHERE medal_type IN ('Gold Medal', 'Silver Medal',
13                          'Bronze Medal')
14 );
```

Execution Plan Window (Edit Data for EXPLAIN (VARCHAR)):

Binary Text

```
1 -> Left hash join (<hash>(mba.discipline)=<hash>(a.Discipline)), extra conditions: (mba.discipline =
2 a.Discipline) (cost=644737605.88 rows=6447353534) (actual time=6.277..111.975 rows=297006 loops=1)
3 -> Nested loop antijoin (cost=2687611.01 rows=26852784) (actual time=3.274..16.645 rows=9069 loops=1)
4 -> Table scan on a (cost=1214.21 rows=11184) (actual time=0.033..4.074 rows=11077 loops=1)
5 -> Single-row index lookup on <subquery3> using <auto_distinct_key> (athlete_name=a.'Name') (actual
6 time=0.001..0.001 rows=0 loops=11077)
7 -> Materialize with deduplication (cost=503.42..503.42 rows=2401) (actual time=10.949..10.949
8 rows=2174 loops=1)
9 -> Filter: (MedalsByAthlete.athlete_name is not null) (cost=263.32 rows=2401) (actual
10 time=0.030..2.220 rows=2401 loops=1)
11 -> Filter: (MedalsByAthlete.medal_type in ('Gold Medal','Silver Medal','Bronze Medal'))
12 (cost=263.32 rows=2401) (actual time=0.028..2.036 rows=2401 loops=1)
13 -> Table scan on MedalsByAthlete (cost=263.32 rows=2401) (actual time=0.025..0.828
14 rows=2401 loops=1)
15 -> Hash
16 -> Table scan on mba (cost=0.02..5.50 rows=240) (actual time=0.001..0.059 rows=735 loops=1)
17 -> Materialize (cost=316.86..322.34 rows=240) (actual time=2.672..2.772 rows=735 loops=1)
18 -> Table scan on <temporary> (cost=0.02..5.50 rows=240) (actual time=0.002..0.081 rows=735
```

Data Length: 1847 bytes

Save... Close

Adding Index on MedalsByAthlete.discipline:

1	-- alter table Coaches add index (NOC(50));
2	• alter table MedalsByAthlete add index (discipline(50));
3	• show indexes from MedalsByAthlete;
4	

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type
MedalsByAthlete	0	PRIMARY	1	athlete_name	A	2174	50	NULL		BTREE
MedalsByAthlete	0	PRIMARY	2	event	A	2401	50	NULL		BTREE
MedalsByAthlete	1	discipline	1	discipline	A	46	50	NULL	YES	BTREE

Runtime:

```

1 • Explain Analyze
2 SELECT a.Name AS LosingAthlete, mba.athlete_name AS WinningAthlete, a.discipline
3 FROM Athletes a
4 LEFT JOIN (
5     SELECT DISTINCT mba.discipline, mba.athlete_name
6     FROM MedalsByAthlete mba
7     WHERE mba.medal_type = 'Gold Medal'
8 ) mba ON a.Discipline = mba.discipline
9 WHERE a.Name NOT IN (
10     SELECT Distinct athlete_name
11     FROM MedalsByAthlete
12     WHERE medal_type IN ('Gold Medal', 'Silver Medal',
13                          'Bronze Medal')
14 )

```


Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

Edit Data for EXPLAIN (VARCHAR)

Binary Text

```

1 -> Limit: 15 row(s) (cost=644737605.88 rows=15) (actual time=8.137..8.151 rows=15 loops=1)
2 -> Left hash join (<hash>(mba.discipline)=<hash>(a.Discipline)), extra conditions: (mba.discipline =
3 a.Discipline) (cost=644737605.88 rows=6447353534) (actual time=8.136..8.149 rows=15 loops=1)
4 -> Nested loop antijoin (cost=2687611.01 rows=26852784) (actual time=4.290..4.294 rows=2 loops=1)
5 -> Table scan on a (cost=1214.21 rows=11184) (actual time=0.045..0.046 rows=2 loops=1)
6 -> Single-row index lookup on <subquery3> using <auto_distinct_key> (athlete_name=a.'Name')
7 (actual time=0.003..0.003 rows=0 loops=2)
8 -> Materialize with deduplication (cost=503.42..503.42 rows=2401) (actual time=4.245..4.245
9 rows=2174 loops=1)
10 -> Filter: (MedalsByAthlete.athlete_name is not null) (cost=263.32 rows=2401) (actual
11 time=0.052..2.924 rows=2401 loops=1)
12 -> Filter: (MedalsByAthlete.medal_type in ('Gold Medal','Silver Medal','Bronze Medal'))
13 (cost=263.32 rows=2401) (actual time=0.050..2.673 rows=2401 loops=1)
14 -> Table scan on MedalsByAthlete (cost=263.32 rows=2401) (actual time=0.046..1.096
15 rows=2401 loops=1)
16 -> Hash
17 -> Table scan on mba (cost=0.02..5.50 rows=240) (actual time=0.002..0.082 rows=735 loops=1)
18 -> Materialize (cost=316.86..322.34 rows=240) (actual time=3.327..3.463 rows=735 loops=1)

```

Data Length: 1979 bytes

Save... Close

Results: Less Efficient, the actual time went from 6.277 - 8.137

Indexing medal_type instead of discipline:

Limit to 1000 rows

```

1 • alter table MedalsByAthlete add index (medal_type(50));
2 • show indexes from MedalsByAthlete;
3

```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type
MedalsByAthlete	0	PRIMARY	1	athlete_name	A	2174	50	NULL		BTREE
MedalsByAthlete	0	PRIMARY	2	event	A	2401	50	NULL		BTREE
MedalsByAthlete	1	medal_type	1	medal_type	A	3	50	NULL	YES	BTREE

Runtime:

```

1 • Explain Analyze
2 SELECT a.Name AS LosingAthlete, mba.athlete_name AS WinningAthlete, a.discipline
3 FROM Athletes a
4 LEFT JOIN (
5     SELECT DISTINCT mba.discipline, mba.athlete_name
6     FROM MedalsByAthlete mba
7     WHERE mba.medal_type = 'Gold Medal'
8 ) mba ON a.Discipline = mba.discipline
9 WHERE a.Name NOT IN (
10     SELECT Distinct athlete_name
11     FROM MedalsByAthlete
12     WHERE medal_type IN ('Gold Medal', 'Silver Medal',
13     'Bronze Medal')
14 )
15 LIMIT 15;

```

Edit Data for EXPLAIN (VARCHAR)

	Binary	Text
1		-> Limit: 15 row(s) (cost=644737605.88 rows=15) (actual time=8.137..8.151 rows=15 loops=1)
2		-> Left hash join (<hash>(mba.discipline)=<hash>(a.Discipline)), extra conditions: (mba.discipline = a.Discipline) (cost=644737605.88 rows=6447353534) (actual time=8.136..8.149 rows=15 loops=1)
3		-> Nested loop antijoin (cost=2687611.01 rows=26852784) (actual time=4.290..4.294 rows=2 loops=1)
4		-> Table scan on a (cost=1214.21 rows=11184) (actual time=0.045..0.046 rows=2 loops=1)
5		-> Single-row index lookup on <subquery3> using <auto_distinct_key> (athlete_name=a.`Name`) (actual time=0.003..0.003 rows=0 loops=2)
6		-> Materialize with deduplication (cost=503.42..503.42 rows=2401) (actual time=4.245..4.245 rows=2174 loops=1)
7		-> Filter: (MedalsByAthlete.athlete_name is not null) (cost=263.32 rows=2401) (actual time=0.052..2.924 rows=2401 loops=1)
8		-> Filter: (MedalsByAthlete.medal_type in ('Gold Medal','Silver Medal','Bronze Medal')) (cost=263.32 rows=2401) (actual time=0.050..2.673 rows=2401 loops=1)
9		-> Table scan on MedalsByAthlete (cost=263.32 rows=2401) (actual time=0.046..1.096 rows=2401 loops=1)
10		-> Hash
11		-> Table scan on mba (cost=0.02..5.50 rows=240) (actual time=0.002..0.082 rows=735 loops=1)
12		-> Materialize (cost=316.86..322.34 rows=240) (actual time=3.327..3.463 rows=735 loops=1)

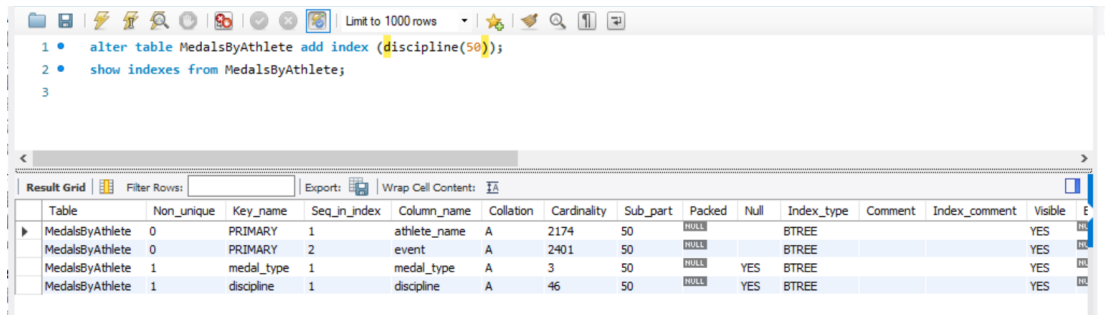
Data Length: 1979 bytes

Save... Close

Results:

Still slower than the default indexing, which is weird since the query has two subqueries that use the 'WHERE clause'

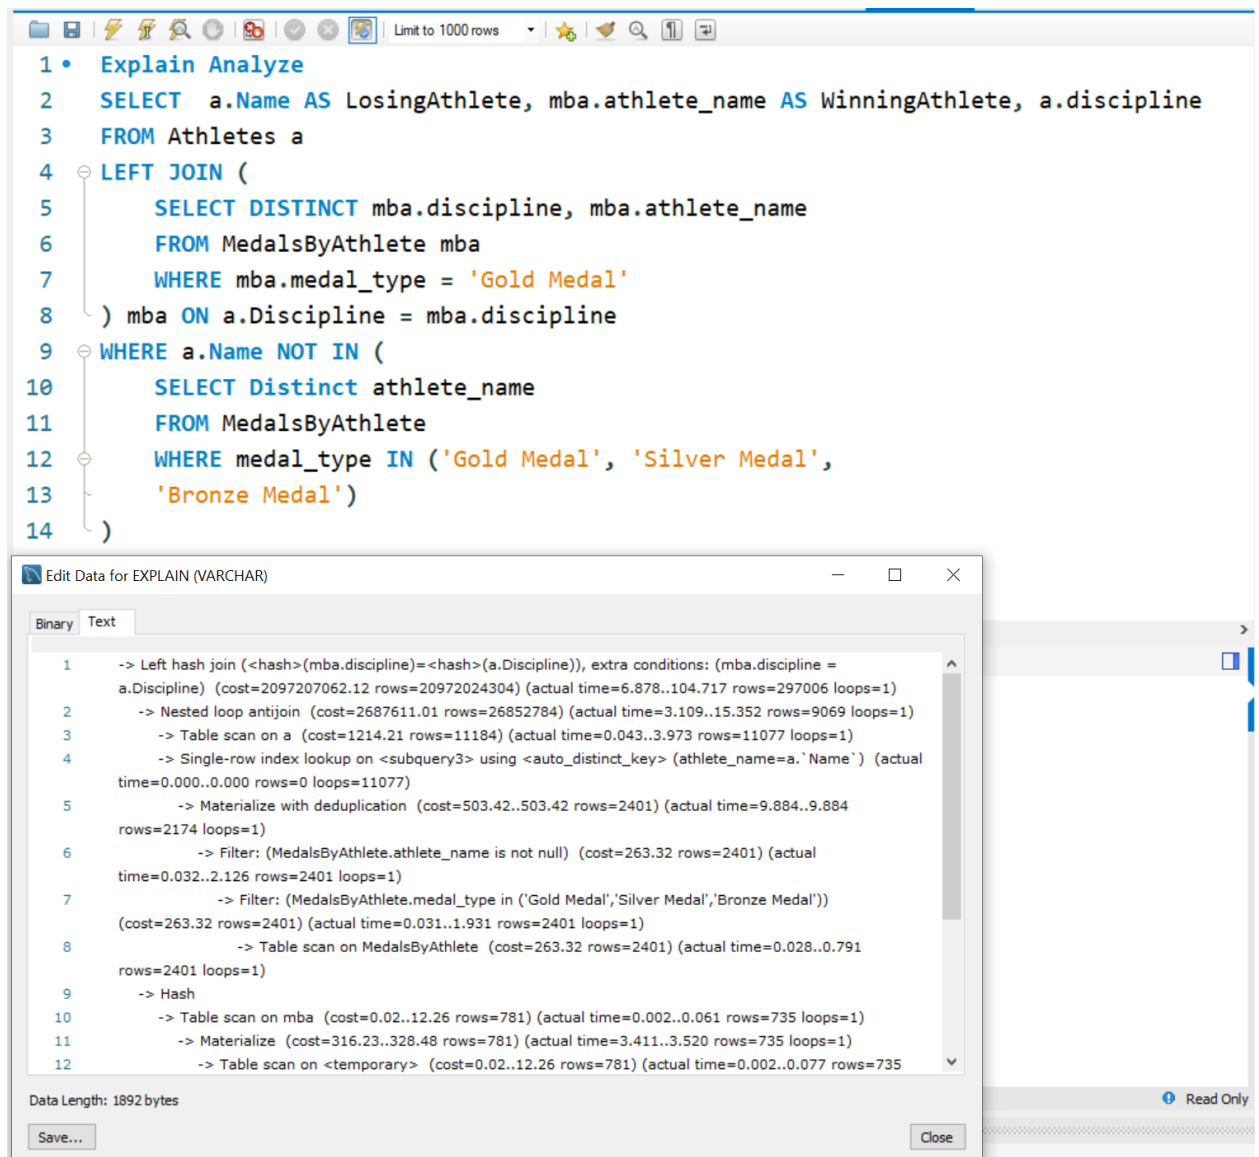
Indexing both discipline and medal_type:



```
1 • alter table MedalsByAthlete add index (discipline(50));
2 • show indexes from MedalsByAthlete;
3
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment	Visible
MedalsByAthlete	0	PRIMARY	1	athlete_name	A	2174	50	HULL		BTREE			YES
MedalsByAthlete	0	PRIMARY	2	event	A	2401	50	HULL		BTREE			YES
MedalsByAthlete	1		1	medal_type	A	3	50	HULL	YES	BTREE			YES
MedalsByAthlete	1		1	discipline	A	46	50	HULL	YES	BTREE			YES

Runtime:



```
1 • Explain Analyze
2 SELECT a.Name AS LosingAthlete, mba.athlete_name AS WinningAthlete, a.discipline
3 FROM Athletes a
4 LEFT JOIN (
5     SELECT DISTINCT mba.discipline, mba.athlete_name
6     FROM MedalsByAthlete mba
7     WHERE mba.medal_type = 'Gold Medal'
8 ) mba ON a.Discipline = mba.discipline
9 WHERE a.Name NOT IN (
10     SELECT Distinct athlete_name
11     FROM MedalsByAthlete
12     WHERE medal_type IN ('Gold Medal', 'Silver Medal',
13                          'Bronze Medal')
14 )
```

Step	Operation	Cost	Rows	Actual Time	Actual Rows	Loops
1	-> Left hash join (<hash>(mba.discipline)=<hash>(a.Discipline)), extra conditions: (mba.discipline = a.Discipline)	2097207062.12	20972024304	6.878..104.717	297006	1
2	-> Nested loop antijoin	2687611.01	26852784	3.109..15.352	9069	1
3	-> Table scan on a	1214.21	11184	0.043..3.973	11077	1
4	-> Single-row index lookup on <subquery3> using <auto_distinct_key> (athlete_name=a.`Name`)	0.000..0.000	0	0.000..0.000	0	11077
5	-> Materialize with deduplication	503.42..503.42	2401	9.884..9.884	2174	1
6	-> Filter: (MedalsByAthlete.athlete_name is not null)	263.32	2401	0.032..2.126	2401	1
7	-> Filter: (MedalsByAthlete.medal_type in ('Gold Medal','Silver Medal','Bronze Medal'))	263.32	2401	0.031..1.931	2401	1
8	-> Table scan on MedalsByAthlete	263.32	2401	0.028..0.791	2401	1
9	-> Hash					
10	-> Table scan on mba	0.02..12.26	781	0.002..0.061	735	1
11	-> Materialize	316.23..328.48	781	3.411..3.520	735	1
12	-> Table scan on <temporary>	0.02..12.26	781	0.002..0.077	735	

Data Length: 1892 bytes

Save... Close

Results:

By using both discipline and medal_type as an index, we were not able to speedup the runtime of the query with both versions costing around 6.8 units of time.

Based on these results, we will stick with default indexing as this produces the shortest runtime.